



RED HAT[®]

AMQ Streams

1.0

Stelios Kousouris
Senior Applications Development Architect

Introduction to Apache Kafka


Apache Kafka ecosystem

- Apache Kafka has an ecosystem consisting of many components / tools
 - Kafka Core
 - Broker
 - Clients library (Producer, Consumer, Admin)
 - Management tools
 - Kafka Connect
 - Kafka Streams
 - Mirror Maker

Apache Kafka ecosystem

Apache Kafka components

- **Kafka Broker**
 - Central component responsible for hosting topics and delivering messages
 - One or more brokers run in a cluster alongside with a Zookeeper ensemble
- **Kafka Producers and Consumers**
 - Java-based clients for sending and receiving messages
- **Kafka Admin tools**
 - Java- and Scala- based tools for managing Kafka brokers
 - Managing topics, ACLs, monitoring etc.

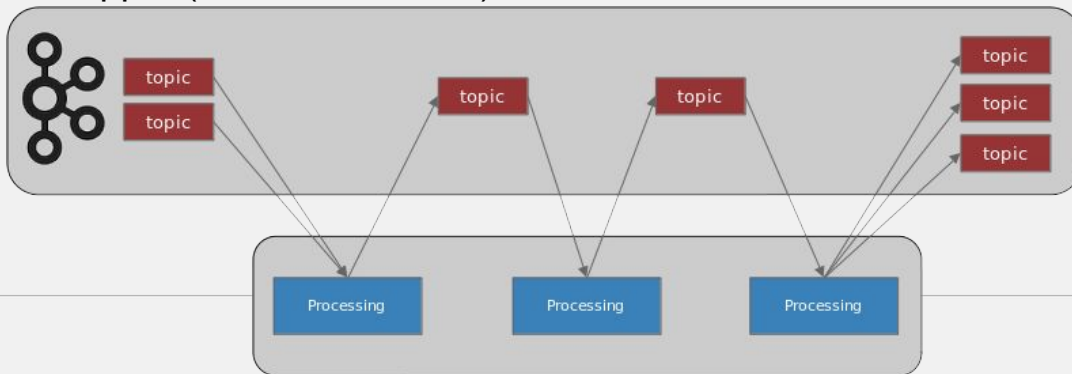


Apache Kafka ecosystem

Apache Kafka components

- **Kafka Streams**

- Stream processing framework
- Streams are Kafka topics (as input and output)
- Scaling the stream application horizontally
- Creates a topology of processing nodes (filter, map, join etc) acting on a stream
- Low level processor API
- High level DSL
- Using “internal” topics (when re-partitioning is needed or for “stateful” transformations)
- Scala wrapper (New in Kafka 2.0)

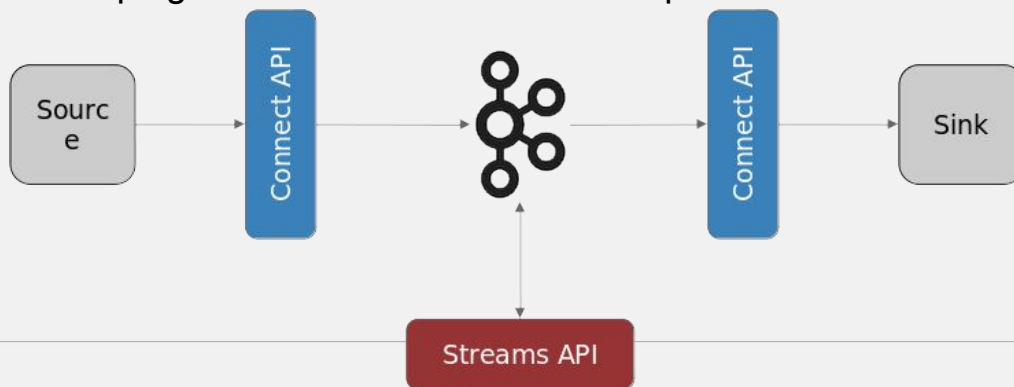


Apache Kafka ecosystem

Apache Kafka components

- **Kafka Connect**

- Framework for transferring data between Kafka and other data systems
- Facilitate data conversion, scaling, load balancing, fault tolerance, ...
- Connector plugins are deployed into Kafka connect cluster
- Well defined API for creating new connectors (with Sink/Source)
- Apache Kafka itself includes only *FileSink* and *FileSource* plugins (reading records from file and posting them as Kafka messages / writing Kafka messages to files)
- Many additional plugins are available outside of Apache Kafka

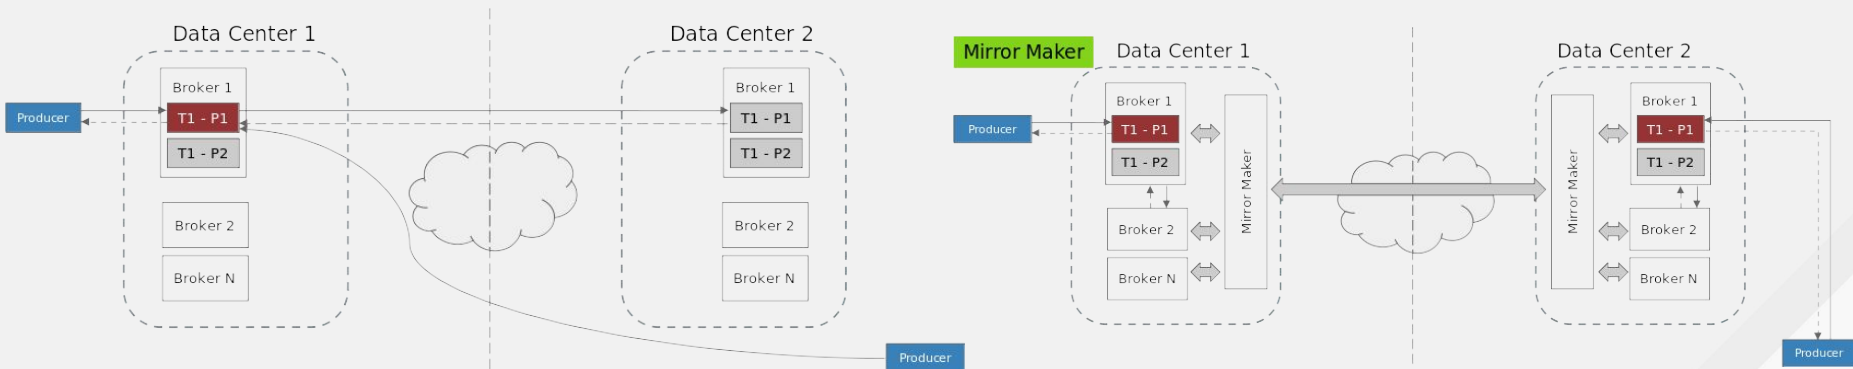


Apache Kafka ecosystem

Apache Kafka components

- **Mirror Maker**

- Kafka clusters do not work well when split across multiple datacenters
- Low bandwidth, High latency
- For use within multiple datacenters it is recommended to setup independent cluster in each data center and mirror the data
- Tool for replication of topics between different clusters




Apache Kafka ecosystem

Apache Kafka components

- **Mirror Maker**

- Kafka clusters do not work well when split across multiple datacenters
- Low bandwidth, High latency
- For use within multiple datacenters it is recommended to setup independent cluster in each data center and mirror the data
- Tool for replication of topics between different clusters



Zookeeper

Apache Kafka internals

- Electing controller
 - Zookeeper is used to elect a controller, make sure there is only one and elect a new one if it crashes
- Cluster membership
 - Which brokers are alive and part of the cluster? This is also managed through ZooKeeper
- Topic configuration
 - Which topics exist, how many partitions each has, where are the replicas, who is the preferred leader, what configuration overrides are set for each topic
- Quotas and ACLs
 - How much data is each client allowed to read and write and who is allowed to do that to which topic

Security

Apache Kafka internals

- Encryption between clients and brokers and between brokers
 - Using SSL
- Authentication of clients (and brokers) connecting to brokers
 - Using SSL (mutual authentication)
 - Using SASL (with PLAIN, Kerberos or SCRAM-SHA as mechanisms)
- Authorization of read/writes operation by clients
 - ACLs on resources such as topics
 - Authenticated “principal” for configuring ACLs
 - Pluggable
- It's possible to mix encryption/no-encryption and authentication/no-authentication

Security

Apache Kafka internals

- Zookeeper
 - No encryption support in latest stable version
 - Support for SASL authentication
 - Using Kerberos or locally stored credentials
- Kafka treats information in public
 - Every user can see it
 - ACL protection can be enabled for changing the data

Supported in 1.0

Supported

Platforms

- AMQ Streams 1.0 server components
 - Broker, Zookeeper, Kafka Connect, MirrorMaker
 - RHEL7 / x86_64 (OpenJDK 8, OracleJDK 8, IBM JDK 8)
- AMQ Streams 1.0 client components
 - Kafka Consumer, Producer and Admin clients
 - Kafka Streams
 - Supported on RHEL7 / x86_64 (OpenJDK 8, OracleJDK 8, IBM JDK 8)

Supported

Protocol & API Versions

- Apache Kafka is backwards compatible with older clients / brokers
 - The protocol / API is versioned according to Kafka release (e.g. 1.1, 1.0, 0.10.1)
 - Latest Kafka still supports many old API versions (from 0.8.0 up to 1.1)
- AMQ Streams 1.0 will support only the latest version
 - The older versions will still work with AMQ Streams 1.0
 - But Red Hat support will be offered only for the latest version
 - In next releases, support for newer versions will be added as they are implemented in Kafka

Supported Components

Broker, MirrorMaker

- Kafka Broker
 - Fully supported on RHEL 7
- MirrorMaker
 - Fully supported

Supported Components

Clients

- Apache Kafka project contains only Java client libraries
 - Consumer, Producer and Admin APIs
 - Streams API
 - Will be supported by AMQ Streams
- Clients for other languages are usually open-source, but are not part of Apache Kafka
 - Many clients are bindings against librdkafka library (C language)
 - These clients will not be supported by AMQ Streams 1.0
 - Users are free to use them on their own

Supported Components

Kafka Connect

- Apache Kafka Connect will be supported in both distributed and standalone modes
- Only the FileSink and FileSource plugins will be supported

Supported Components

Apache Zookeeper

- Apache Zookeeper is a dependency of Kafka
- Just an implementation detail for AMQ Streams
- Zookeeper will be supported only for use by Kafka
 - No support will be provided for using Zookeeper with other applications
 - Plans for removing Zookeeper dependency are being discussed upstream
 - Once Zookeeper is not needed by Apache Kafka, we will drop it as well

Unsupported

- Several notable parts of wider Kafka ecosystem which will not be supported in 1.0
 - Non-Java clients
 - Kafka Connect plugins
 - Confluent REST proxy (<https://github.com/confluentinc/kafka-rest>)
 - Confluent Schema Registry (<https://github.com/confluentinc/schema-registry>)
 - KSQL (<https://github.com/confluentinc/ksql>)

After 1.0 - Looking to support

- Security
 - Improve authentication and authorization options
 - Support LDAP for authentication
 - Support Red Hat SSO for authentication and authorization
 - Support authentication and authorization using tokens
- AMQP Bridge
 - Based on <https://github.com/strimzi/amqp-kafka-bridge>
 - Advantages
 - Integrate with other AMQ products using AMQP
 - Use our supported AMQP clients (as an alternative for non-Java Kafka clients which we don't support)
 - No need to expose the whole Kafka cluster
 - Further plans to extend the Bridge with HTTP and MQTT support

After 1.0 - Looking to support

- Schema Registry
 - Possibly based on <https://github.com/jhalliday/perspicuus>
 - Advantages
 - Support for different schemas
 - JSON, AVRO, Protocol Buffers
- Cluster balancing
 - Support automated cluster balancer
 - Automatically balance the cluster
 - Move topics between nodes to achieve best performance
 - Based on Memory, CPU, Disk or Network IO
 - Reusing one of existing projects
 - Build from scratch using OptaPlanner for optimization

Why & Where should you use AMQ Streams

Why use AMQ Streams

- Scalability and Performance
 - Designed for horizontal scalability
 - Cluster sizes from few brokers up to 1000s of brokers
 - 3 nodes usually seen as minimum for production (HA, message durability)
 - Most clusters are under 50 nodes
 - Different approaches: One big cluster vs. several small clusters
 - Scaling has minimal impact on throughput and latency
 - Adding nodes to running cluster is easy
- Message ordering guarantee
 - Messages are written to disk in the same order as received by the broker
 - Messages are read from disk from the requested offset
 - Kafka protocol makes sure that one consumer can read only one partition
 - Note: Order is guaranteed only within a single partition!
 - Note: No synchronization between producers!

Why use AMQ Streams

- Message rewind / replay
 - Limited only by available disk space
 - Amount of stored messages has no impact on performance
 - Topic / Partition size has no direct impact on performance
 - Allows to reconstruct application state by replaying the messages
 - Combined with compacted topics allows to use Kafka as key-value store
 - Event sourcing (<https://martinfowler.com/eaaDev/EventSourcing.html>)
 - Parallel running (https://en.wikipedia.org/wiki/Parallel_running)

AMQ Streams

vs AMQ 7 Broker

- Use AMQ Streams when
 - You need scalability to achieve highest possible throughput
 - You need large number of parallel consumers
 - Your messages are cattle and not pets
 - You need message ordering or replay features
 - You can reuse some of the components available in Kafka ecosystem instead (e.g. Kafka Connect plugins)

AMQ Streams

vs AMQ 7 Broker

- Use AMQ Broker when
 - You care about individual messages
 - You use request/response patterns
 - You want to use standard clients and protocols
 - You need TTL or DLQ semantics
 - You need more sophisticated routing patterns

AMQ Streams on Openshift

AMQ Streams on OpenShift

What is it?

- Based on OSS project called Strimzi
- Provides:
 - Docker images for running Apache Kafka and Zookeeper
 - Tooling for managing and configuring Apache Kafka clusters and topics
- Follows the Kubernetes operator model
- OpenShift 3.9 and higher

AMQ Streams on OpenShift

What is Strimzi?

- Open source project focused on running Apache Kafka on Kubernetes and OpenShift
 - Web site: <http://strimzi.io/>

AMQ Streams on OpenShift

Why is Strimzi?

- Kafka brokers inherently stateful.
- Updating and scaling a Kafka cluster requires careful orchestration to ensure that messaging clients are unaffected and no records are lost.
- By design, Kafka clients connect to all the brokers in the cluster.
- This is part of what gives Kafka its horizontal scaling and high availability, but when running on OpenShift, this means the Kafka cluster cannot simply be put behind a load-balanced service like other services. Instead services have to be orchestrated in parallel with cluster scaling.
- Running Kafka also requires running a Zookeeper cluster, which has many of the same challenges as running the Kafka cluster.

AMQ Streams on OpenShift - How?

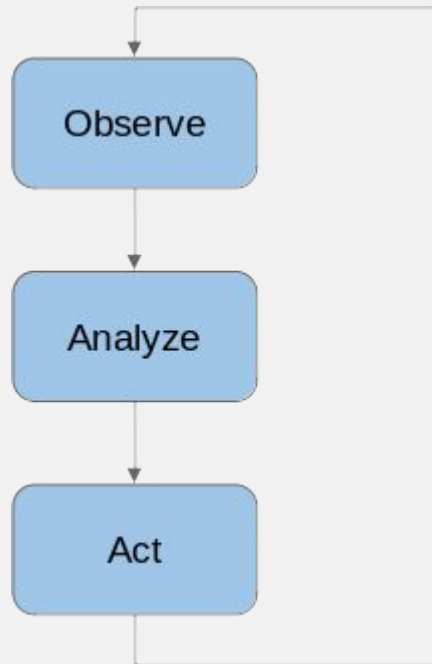
Operator

- Application-specific controller is used to create, configure and manage other complex application
 - The controller contains specific domain / application knowledge
 - Usually used for stateful applications (databases, ...) which are non-trivial to operate on Kubernetes / OpenShift
- Controller operates based on input from ~~Config Maps~~ or **Custom Resource Definitions**
 - User describes the desired state
 - Controller applies this state to the application

AMQ Streams on OpenShift - How?

Operator

- Observe
 - Monitor the current state of the application
- Analyze
 - Compare the actual state to the desired state
- Act
 - Resolve any differences between actual and desired state



AMQ Streams on OpenShift - How?

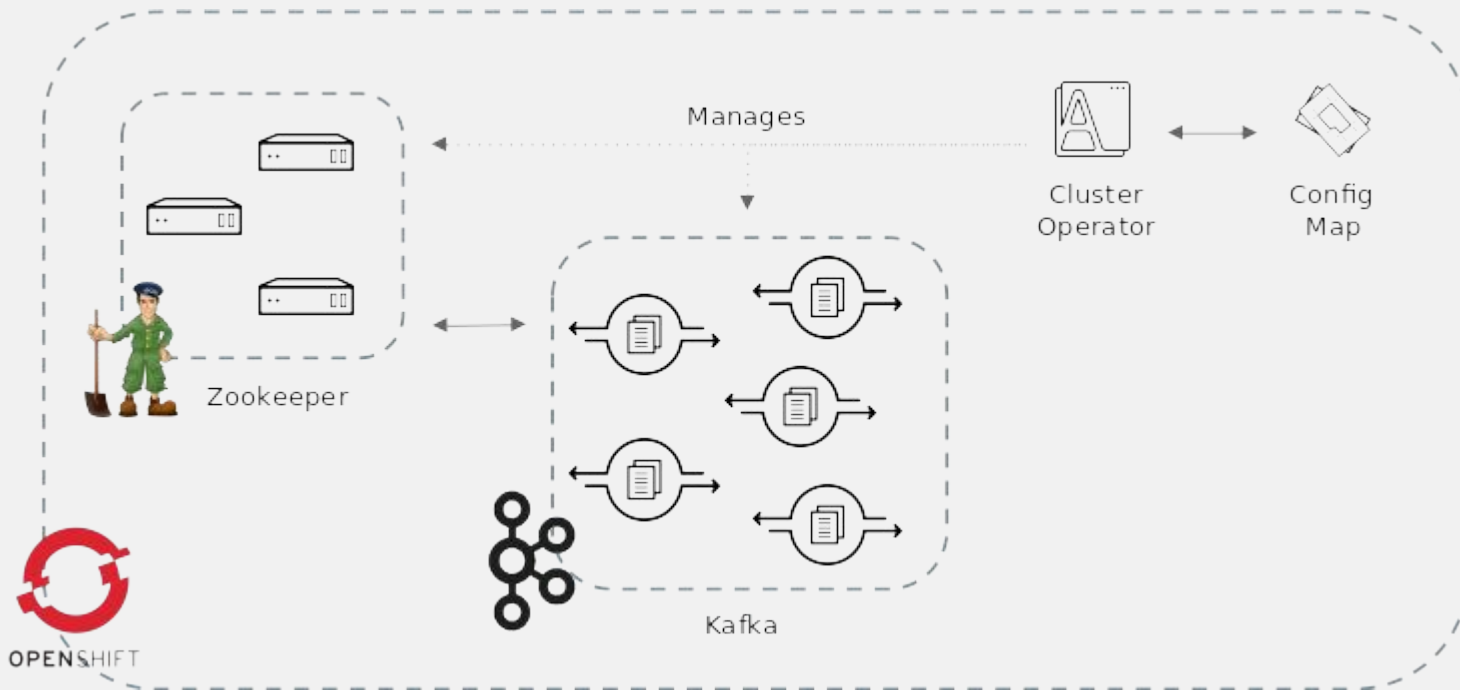
Custom Resource Definitions

- Flexible data structure
- Possibility to set permissions for the CRD resources

Cluster Operator

Cluster Operator

Creating and managing Apache Kafka clusters



Cluster Operator

Supported clusters

- Cluster Operator can deploy and manage two kinds of cluster
 - Kafka cluster
 - Cluster of Kafka brokers
 - Includes Zookeeper deployment
 - Using Stateful Sets for managing Kafka and Zookeeper
- Kafka Connect cluster
 - Distributed Kafka Connect cluster
 - Using Deployment
 - S2I support for adding additional plugins

Cluster Operator

Supported features

- Cluster Operator currently allows to configure
 - Number of Zookeeper, Kafka and Kafka Connect nodes
 - Configuration of Kafka and Kafka Connect
 - Storage
 - Persistent versus Ephemeral
 - Storage size
 - Metrics exports for Prometheus
 - Healthchecks

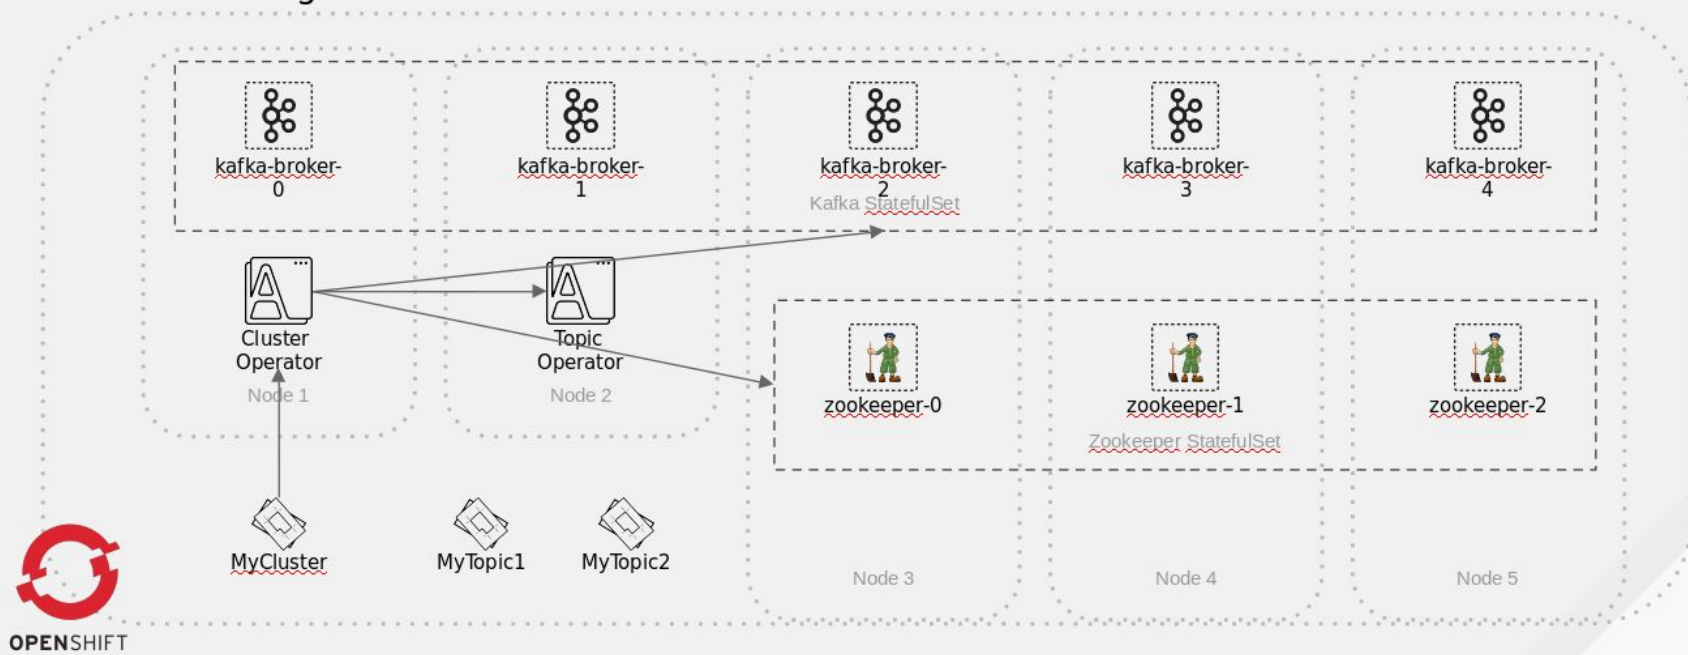
Cluster Operator

Creating a cluster

- One Cluster Operator can manage several clusters in parallel
 - Can cover one or more projects / namespaces
- To deploy new cluster
 - (Deploy the cluster controller)
 - Create a Config Map describing the cluster
 - The Controller will see the Config Map and start deploying the cluster
 - The cluster will be deployed and ready to use

Cluster Operator

Creating a cluster



Cluster Operator

Managing cluster

- Clusters can be modified by modifying the Config Map
 - Scale-up / scale-down
 - Kafka configuration
- Cluster Operator will update the cluster to match the desired state described in Config Map
- Update does not allow to change storage configuration
 - Such operation cannot be done without losing data

Cluster Operator

Deleting cluster

- Clusters can be deleted by deleting the CRDs
 - All cluster resources will be removed
 - Persistent Volume Claims will be deleted according to user configuration
- Deleting the CRDs is irreversible

Cluster Operator

Kafka Connect and Source 2 Image

- Kafka Connect can be deployed with Source 2 Image support (S2I)
 - By default Kafka Connect contains only FileSync and FileSource plugins which are not much useful in distributed mode
 - To make it more useful, additional plugins need to be added
 - New Docker image with additional plugins can be build using S2I
 - User prepares a binary of the plugin and triggers new build
 - The build adds the plugin to the original Kafka Connect image
 - Rolling update is triggered to use the new image

Cluster Operator

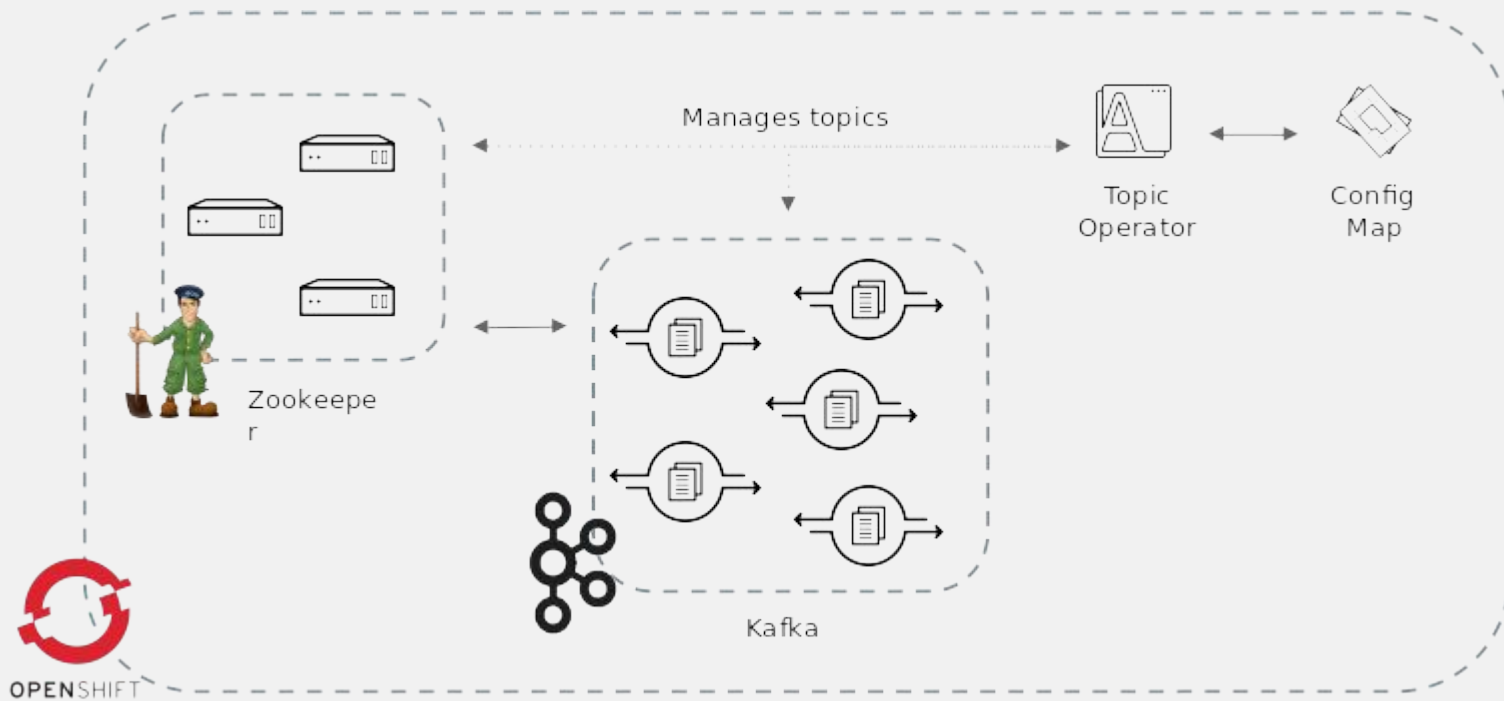
Kafka Connect and Source 2 Image

- Kafka Connect can be deployed with Source 2 Image support (S2I)
 - By default Kafka Connect contains only FileSync and FileSource plugins which are not much useful in distributed mode
 - To make it more useful, additional plugins need to be added
 - New Docker image with additional plugins can be build using S2I
 - User prepares a binary of the plugin and triggers new build
 - The build adds the plugin to the original Kafka Connect image
 - Rolling update is triggered to use the new image

Topic Operator

Topic Operator

Creating and managing Kafka topics



Topic Operator

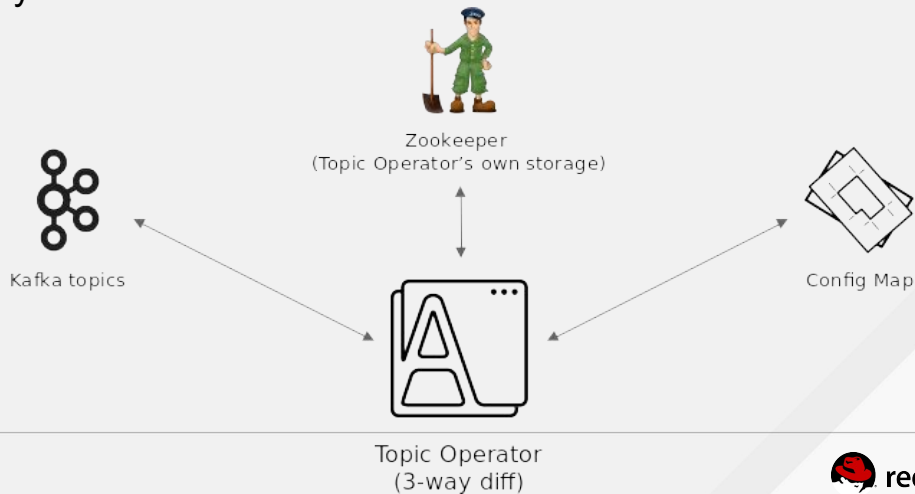
Creating and managing Kafka topics

- Topic Operator is managing Kafka topics
- Some Kafka components (Streams, Connect) often create their own topics
 - Bi-directional synchronization
 - Changes done directly in Kafka / Zookeeper are applied to CRDs
 - Changes done in Config Maps are applied to Kafka topics
- Topic Operator solves this by using 3-way diff
 - Our own Zookeeper-based store
 - Apache Kafka / Zookeeper
 - CRDs

Topic Operator

Creating and managing Kafka topics

- Topic Operator is managing Kafka topics
- Some Kafka components (Streams, Connect) often create their own topics
 - Bi-directional synchronization
 - Changes done directly in Kafka / Zookeeper are applied to CRDs
 - Changes done in Config Maps are applied to Kafka topics
- Topic Operator solves this by using 3-way diff
 - Our own Zookeeper-based store
 - Apache Kafka / Zookeeper
 - CRDs



Topic Operator

Creating and managing Kafka topics

- Kafka gives user more freedom for naming than OpenShift
 - CRDs configuring topics contain “name” field to define the name which is not allowed for OpenShift resource
 - If such topic is created by Kafka, it will be mapped to a Config Map named after specially encoded name
- We recommend to use only topic names which are allowed as OpenShift resource names

User Operator

User Operator

Creating and managing Kafka users/acls

- The User Operator provides synchronization between a *KafkaUser* custom resource in OpenShift and Kafka's own user and Access Control List data structures.
- Allows user to provision the user accounts their application needs at the same time, and in the same way as the application itself: As OpenShift CRD resources.
- User Operator watches for *KafkaUser* custom resource creation and then ensures they are configured in the KAFKA cluster.
- It does not sync changes from KAFKA cluster to Openshift CRDs (unlike Topic Operator)
- In addition manages authorization rule by including description of the user's rights on specific topics

User Operator

```
apiVersion: kafka.strimzi.io/v1alpha1
kind: KafkaUser
metadata:
  name: my-user
  labels:
    strimzi.io/cluster: my-cluster
spec:
  authentication:
    type: tls
  authorization:
    type: simple
    acls:
      - resource:
          type: topic
          name: my-topic
          operation: Read
      - resource:
          type: topic
          name: my-topic
          operation: Describe
      - resource:
          type: group
          name: my-hello-world-consumer
          operation: Read
      - resource:
          type: topic
          name: my-topic
          operation: Write
      - resource:
          type: topic
```



THANK YOU



plus.google.com/+RedHat



facebook.com/redhatinc



linkedin.com/company/red-hat



twitter.com/RedHatNews



youtube.com/user/RedHatVideos

